





Jailbreaking defenses

Not just topics but also **alignment-relevant** behavior



Separability peaks in

Source: (LM semantics, Saglam et al., 2025).



Consistent clustering in the final layer across 5 chat models on the WildJailbreak dataset

Jailbreaking defenses

Defense from Within: A Simple Latent-Space Guardrail

Train two linear (e.g., SVM), one nonlinear (e.g., MLP) classifiers:

1. Direct vs. Indirect: f_1^{SVM}
2. (Direct) Benign vs. Malicious: f_2^{SVM}
3. (Indirect) Benign vs. Malicious: f_3^{MLP}

