

Agenda: Tutorial on jailbreaking LLMs

► Part I: **Attacks**

- *You'll learn* how to **trick LLMs** into generating harmful text.

► Part II: Defenses

- *You'll learn* how to **secure LLMs** against malicious attacks.

► Part III: **Agents**

- *You'll learn* about the **emerging risks** of deploying AI agents.

Hamed Hassani [40 minutes]

AlexRobbey [40 minutes]

Amin Karbasi [30 minutes]

Agenda: Tutorial on **jailbreaking** LLMs

- ▶ Part I: **Attacks** **Hamed Hassani** [40 minutes]
 - ▶ *You'll learn* how to **trick LLMs** into generating harmful text.
- ▶ Part II: **Defenses** **Amin Karbasi** [30 minutes]
 - ▶ *You'll learn* how to **secure LLMs** against malicious attacks.
- ▶ Part III: **Agents** **Alex Robey** [40 minutes]
 - ▶ *You'll learn* about the **emerging risks** of deploying AI agents.

Agenda: Tutorial on **jailbreaking** LLMs

- ▶ What we *won't* cover today
 - ▶ Hallucination
 - ▶ Privacy and memorization
 - ▶ Attacks on modalities other than text generation
 - ▶ Model weight exfiltration
 - ▶ Alignment faking, scheming, and AI control

