

The capacity of a LLM-enabled system to **perceive its environment**,
maintain internal state, and **autonomously choose actions** that
causally influence the external world.

What is an *agent*?

What are the risks of deploying **agents**?

What are the risks of deploying **agents**?

What are the risks of deploying **agents**?

