



Open directions

There is no perfect method yet; lots of opportunities for future research.



Is bigger always safer?

(Scaling laws; Panfilov et al., 2025).



Is post-hoc mitigation required for safety?

(STACK; McKensize et al., 2025).



Are jailbroken outputs useful?

(Jailbreak tax; Nikolic et al., 2025).

Agenda: Tutorial on **jailbreaking** LLMs

- ▶ Part I: **Attacks** **Hamed Hassani** [40 minutes]
 - ▶ *You'll learn* how to **trick LLMs** into generating harmful text.
- ▶ Part II: **Defenses** **Amin Karbasi** [30 minutes]
 - ▶ *You'll learn* how to **secure LLMs** against malicious attacks.
- ▶ Part III: **Agents** **Alex Robey** [40 minutes]
 - ▶ *You'll learn* about the **emerging risks** of deploying AI agents.