

Agenda: Tutorial on **jailbreaking** LLMs

- ▶ What we *won't* cover today
 - ▶ Hallucination
 - ▶ Privacy and memorization
 - ▶ Attacks on modalities other than text generation
 - ▶ Model weight exfiltration
 - ▶ Alignment faking, scheming, and AI control



Agenda: Tutorial on **jailbreaking** LLMs

▶ Part I: **Attacks**

Hamed Hassani [40 minutes]

- ▶ *You'll learn* how to **trick LLMs** into generating harmful text.

▶ Part II: **Defenses**

Amin Karbasi [30 minutes]

- ▶ *You'll learn* how to **secure LLMs** against malicious attacks.

▶ Part III: **Agents**

Alex Robey [40 minutes]

- ▶ *You'll learn* about the **emerging risks** of deploying AI agents.