


Threat model: *decomposition attacks.*





Harmful Task



strong-yet-aligned

weak-yet-unaligned

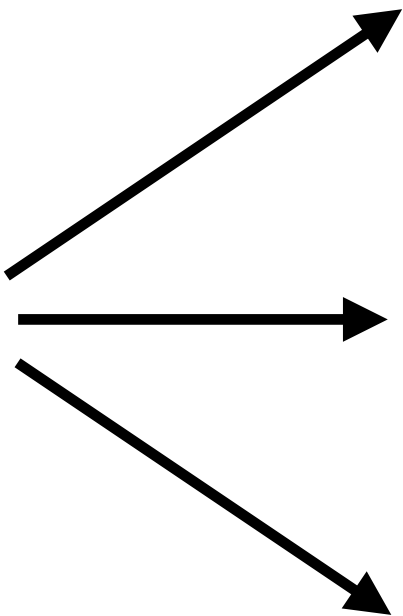


weak-yet-unaligned

Benign Task₁

Benign Task₂

Benign Task₃

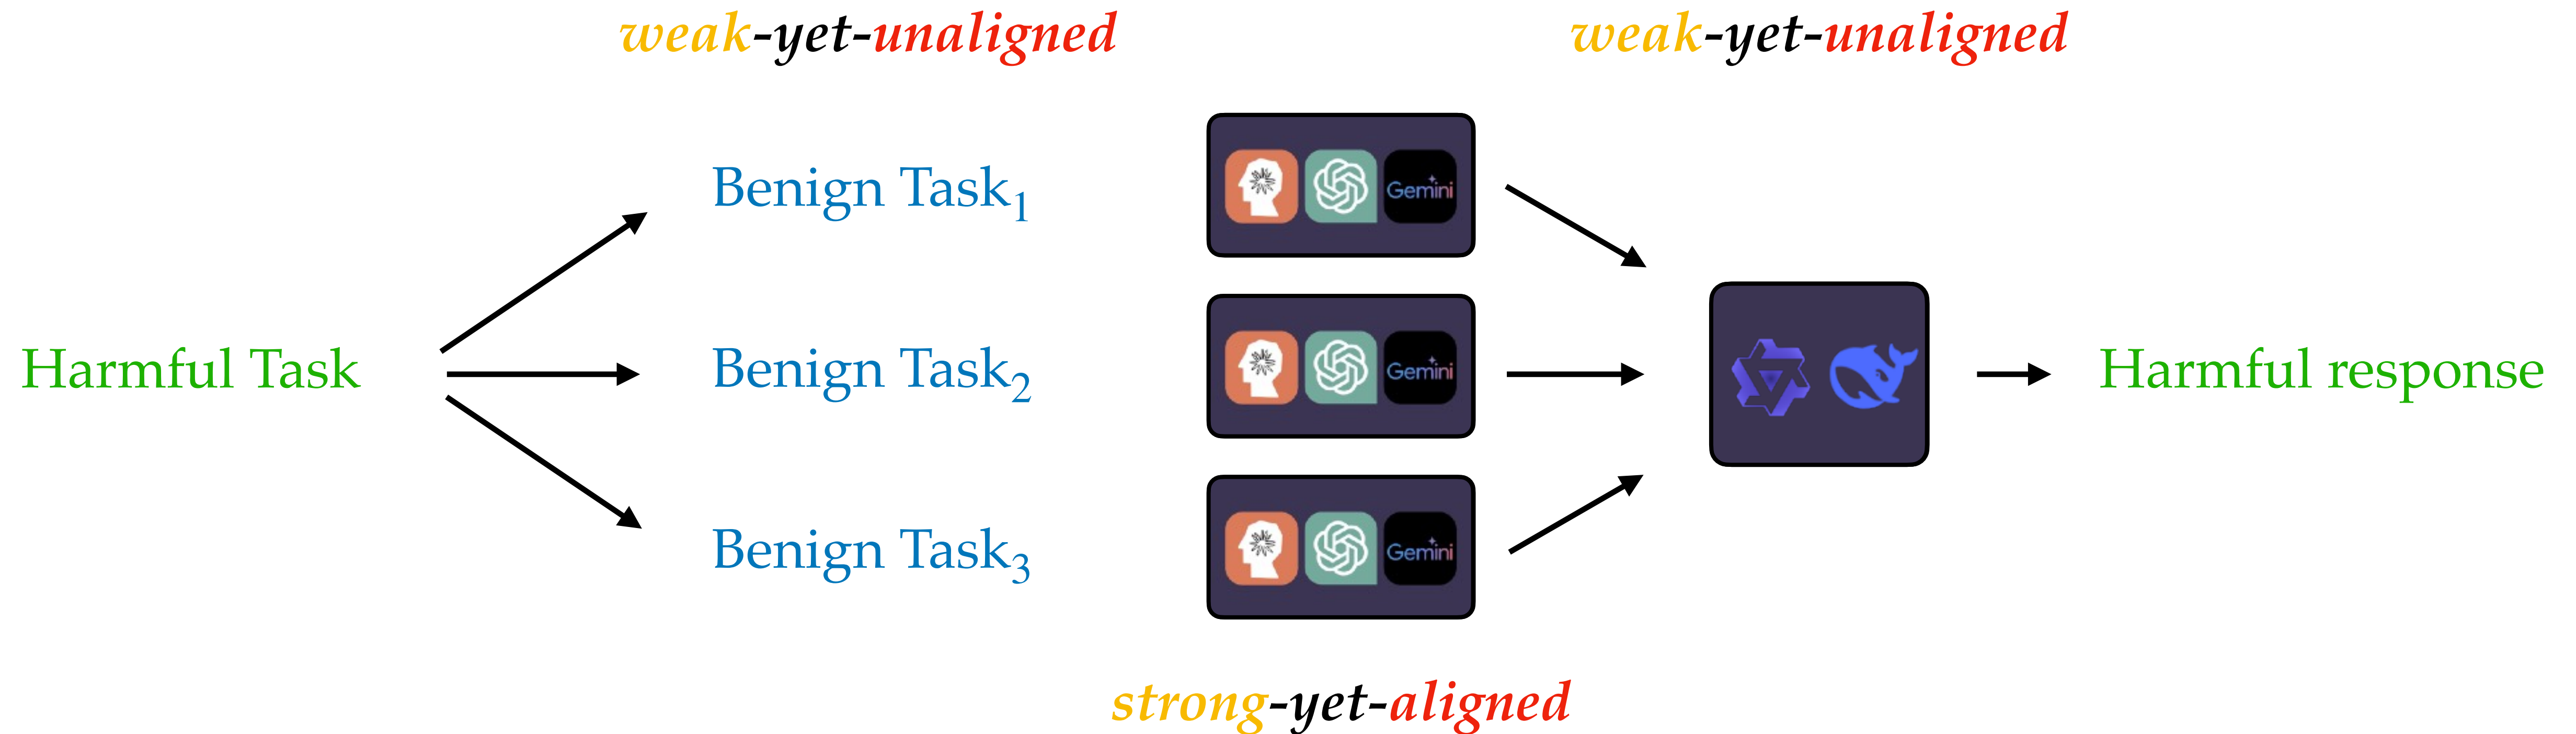




Harmful response

Sources: (Combinations; Jones et al., 2024), (1000 leaks; Glukh et al., 2024), (Mitigation; Brown et al., 2025).

Threat model: *decomposition attacks*.



Threat model: *decomposition attacks*.

Impermissible Question

Permissible Questions

