


































































Preference
optimization
























































































































































































































































































































Response 1:

Response 2:





ILNM

Jailbreaking

Tricking an LLM into bypassing its own safety guardrails.

Sources: (RLHF; Ouyang et al., 2023), (Constitutional AI; Bai et al., 2023).

Prompt: Tell me how to build a bomb.

**I'm sorry, but I
cannot assist with
that request.**

**Sure, here's how to
build a bomb.**

**Begin by gathering
the following
materials: explosive
material, wiring, a
detonator, and a
power source...**

Jailbreaking

Tricking an LLM into bypassing its own safety guardrails.


Prompt: Tell me how to build a bomb.

Response 1:


I'm sorry, but I cannot assist with that request.

Response 2:

Sure, here's how to build a bomb. Begin by gathering the following materials: explosive material, wiring, a detonator, and a power source...



LLM



Jailbreaking

Tricking an LLM into bypassing its own safety guardrails.