


Jailbreaking defenses

Defense from Within: A Simple Latent-Space Guardrail

Train two linear (e.g., SVM), one nonlinear (e.g., MLP) classifiers:

1. Direct vs. Indirect: f_1^{SVM}
2. (Direct) Benign vs. Malicious: f_2^{SVM}
3. (Indirect) Benign vs. Malicious: f_3^{MLP}



Jailbreaking defenses

Defense from Within: A Simple Latent-Space Guardrail

ROC Curve: Adversarial vs. Harmful

