


arXiv abstracts
across topics

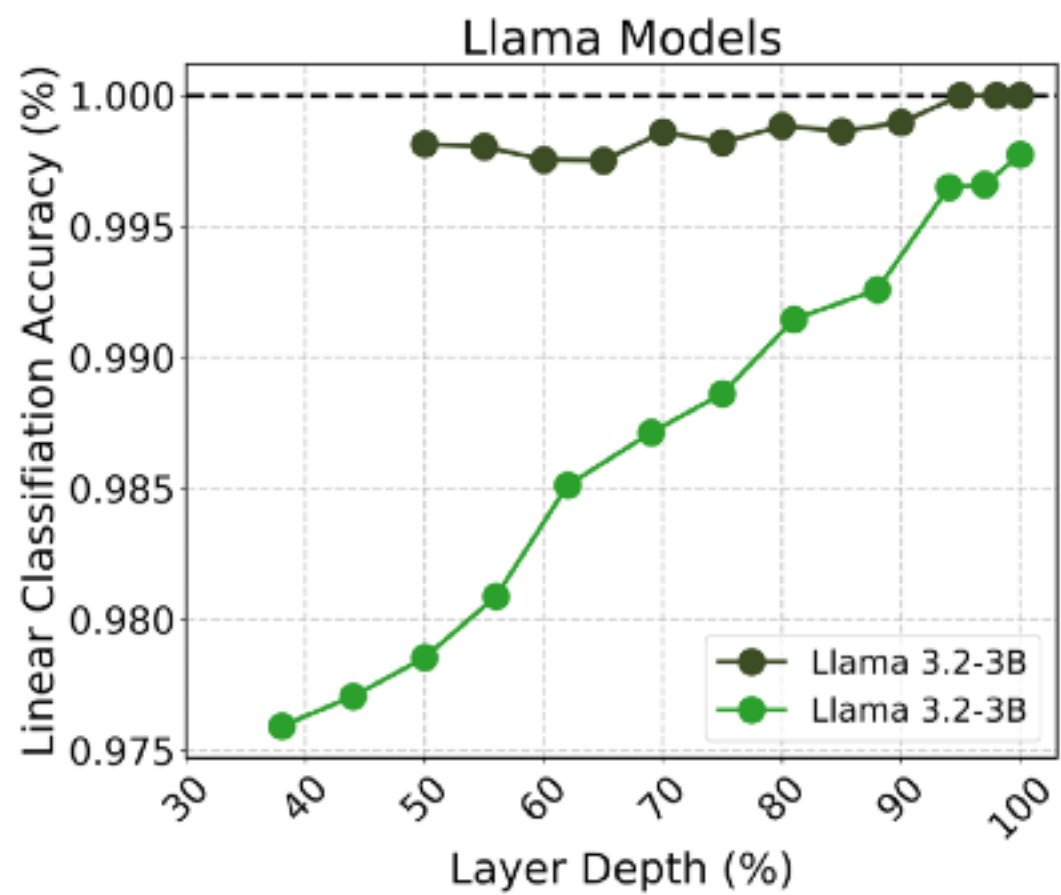
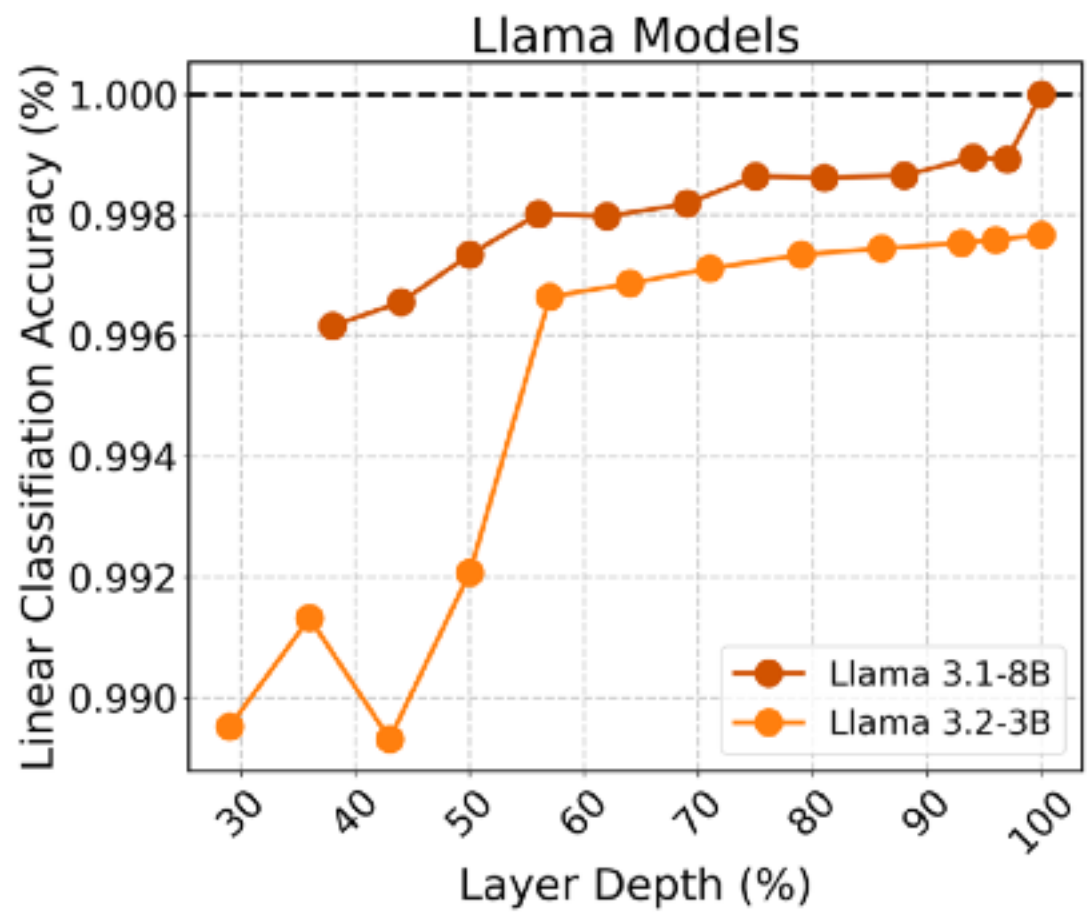
Jailbreaking defenses

Linearly separable clusters in $\mathbb{R}^d$!



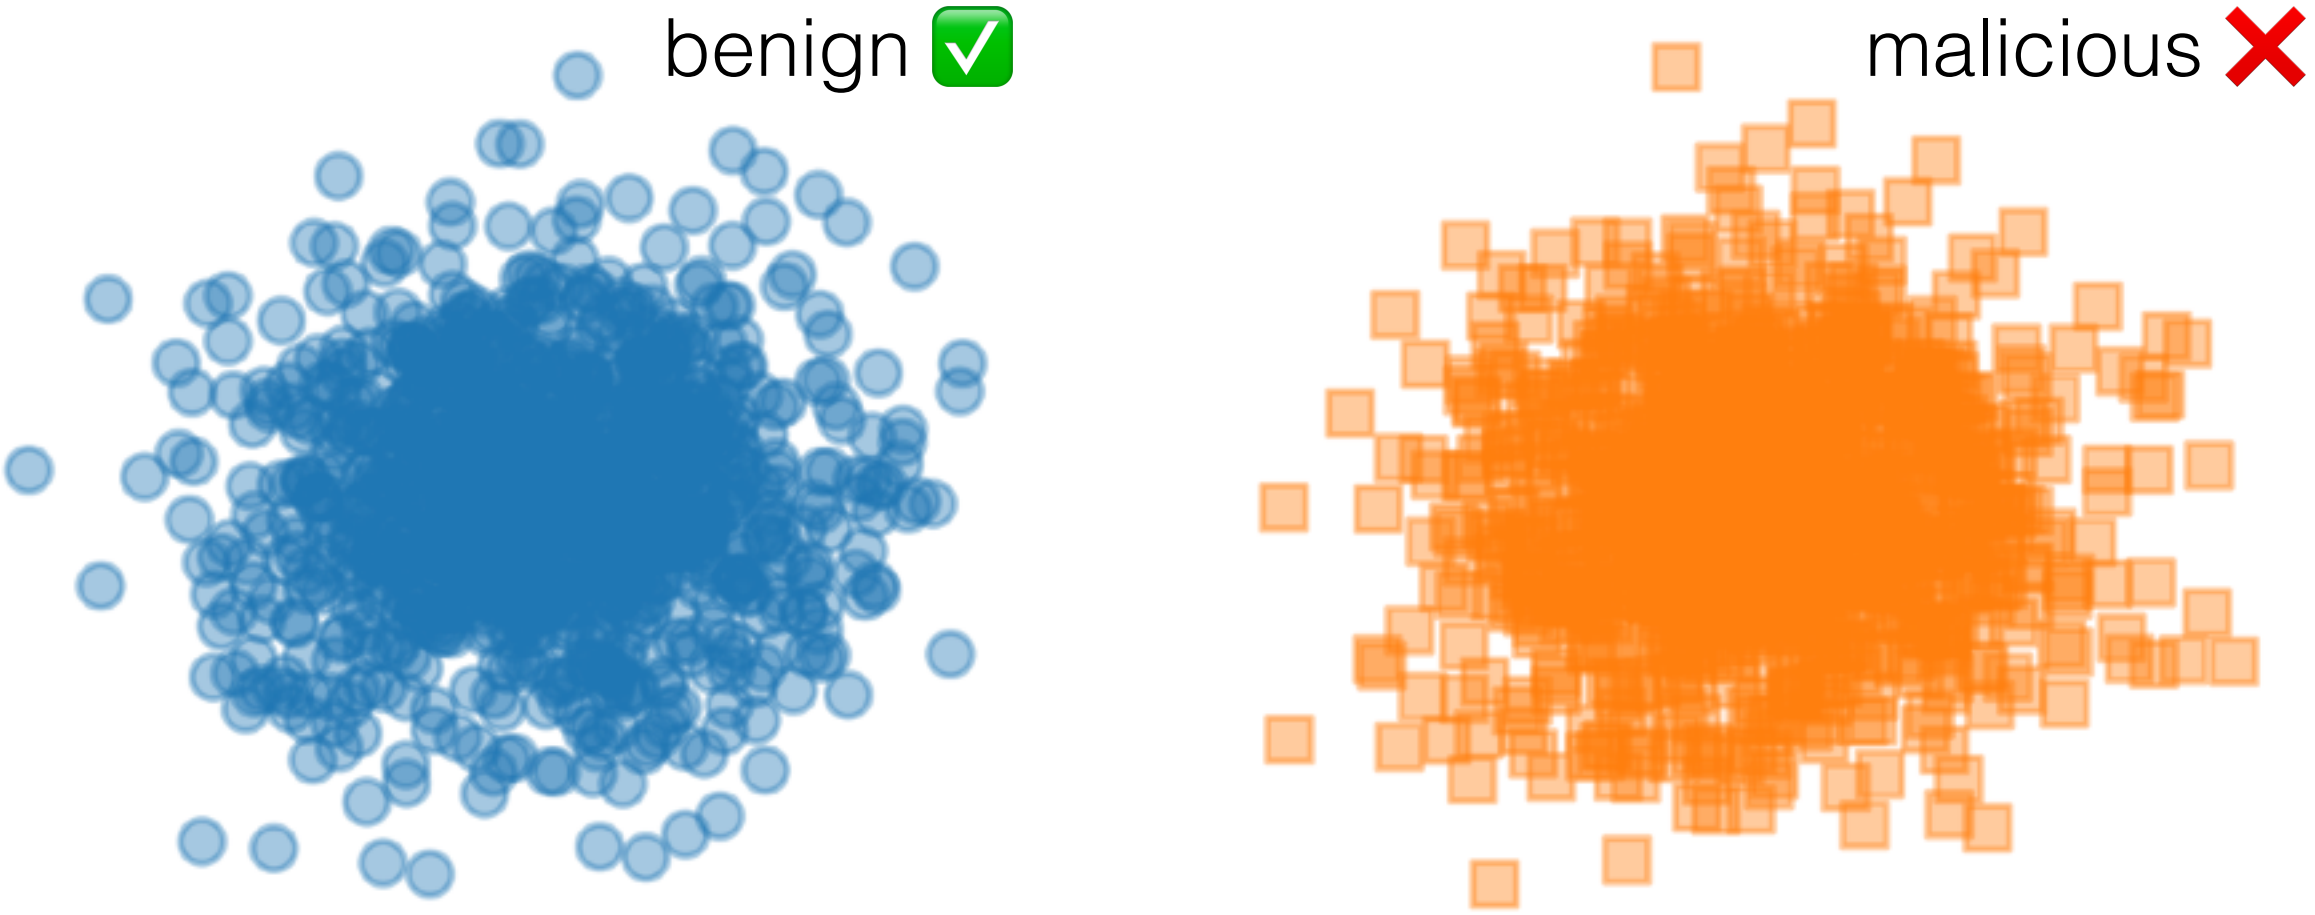
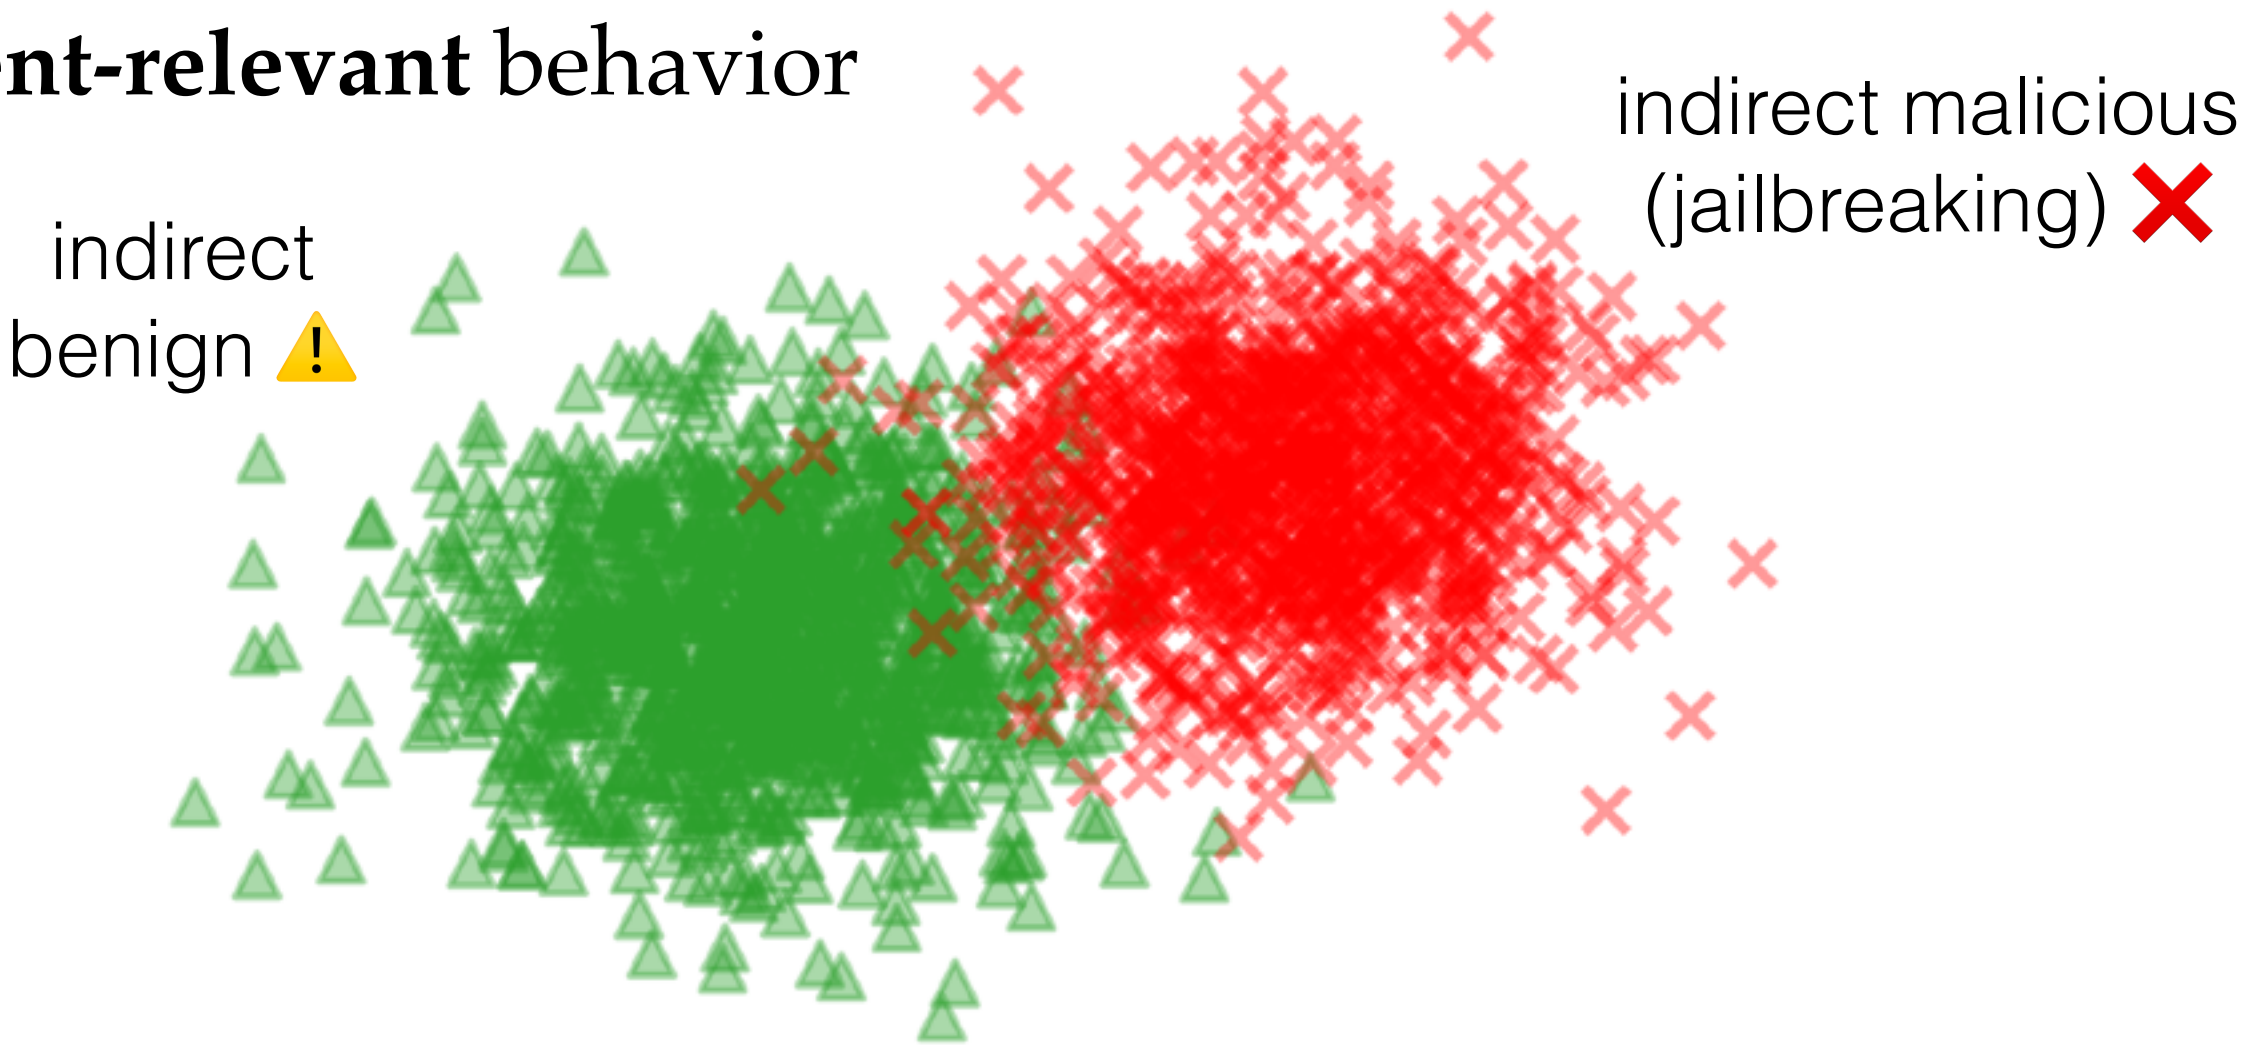
Jailbreaking defenses

Not just topics but also **alignment-relevant** behavior



Separability peaks in

Source: (LM semantics, Saglam et al., 2025).



Consistent clustering in the final layer across 5 chat models on the WildJailbreak dataset