

Conclusion: Open problems & future directions

You're now an expert on jailbreaking LLMs and agents!

► Open problems

► We have powerful attacks; we lack general-purpose defenses, especially for multi-turn and decomposition attacks

► Defenses for agents, possibly domain specific; constructing an instruction hierarchy that accounts for capabilities, external tools, world models, etc.

► Look beyond textual models: robotics, multi-modal, etc.

► Risks of fully autonomous AI agents

Conclusion: Open problems & future directions

- ▶ You're now an expert on jailbreaking LLMs and agents!
- ▶ Open problems
 - ▶ We have powerful attacks; we lack general-purpose defenses, especially for multi-turn and decomposition attacks
 - ▶ Defenses for agents, possibly domain specific; constructing an instruction hierarchy that accounts for capabilities, external tools, world models, etc.
 - ▶ Look beyond textual models: robotics, multi-modal, etc.
 - ▶ Risks of fully autonomous AI agents

Jailbreaking LLMs and Agentic Systems:

Attacks, **Defenses**, and **Evaluations**

ICML 2025

Any questions?