


StrongREJECT scaling (compute coefficient)

o1-mini

o1-preview

o1



CoT-based safety also facilitates trading inference-time compute for adversarial robustness. I.e., safer model when scaling the test-time compute [OpenAI, 2025]:



Safer
Answers


Deliberative alignment

Jailbreaking defenses

Jailbreaking defenses

Deliberative alignment

CoT-based safety also facilitates trading inference-time compute for adversarial robustness. I.e., safer model when scaling the test-time compute [OpenAI, 2025]:



Jailbreaking defenses

Deliberative alignment

1. **Attack mitigation.** Robustness against adaptive attacks. 😊
2. **Non-conservatism.** Maintain the ability to generate realistic text. 😊
3. **Efficiency.** Avoid retraining, maximize query efficiency. 😐
4. **Compatibility.** White- & black-box attacks. 😐