


Jailbreaking defenses

Defense from Within: A Simple Latent-Space Guardrail

ROC Curve: Adversarial vs. Harmful



94.37%

OVERALL ACCURACY

79.30%

ADVERSARIAL MACRO F1

0.9912

AUC

0.95

WEIGHTED F1

Confusion Matrix

	0	1	2	3
0	1000	0	0	0
1	0	1000	0	0
2	0	1	168	41
3	0	54	141	1805
Class 0: Vanilla Benign	Class 1: Vanilla Harmful	Class 2: Adversarial Benign	Class 3: Adversarial Harmful	

Performance Metrics

CLASS	PRECISION	RECALL	F1- SCORE	SUPPORT
Vanilla Benign	1.00	1.00	1.00	1000
Vanilla Harmful	0.95	1.00	0.97	1000
Adversarial Benign	0.54	0.80	0.65	210
Adversarial Harmful	0.98	0.90	0.94	2000

Jailbreaking defenses

Guardrail

1. **Attack mitigation.** Robustness against adaptive attacks. 
2. **Non-conservatism.** Maintain the ability to generate realistic text. 
3. **Efficiency.** Avoid retraining, maximize query efficiency. 
4. **Compatibility.** White- & black-box attacks. 